LLMs in Drug Discovery: Methods, Evaluation, Ethics, and Future Directions

Abstract

This review explores the burgeoning field of Large Language Models (LLMs) in drug discovery, focusing on novel methodologies, evaluation metrics, and ethical considerations. It synthesizes findings from studies employing LLMs for tasks like de novo molecule generation, property prediction, and target identification, with a particular emphasis on generative models such as CharRNN, VAE, and JTN-VAE. A key challenge identified is the lack of standardized evaluation and benchmarking protocols, hindering fair comparisons between different LLM architectures and methodologies. The review highlights the importance of platforms like MOSES in addressing this issue. Furthermore, it delves into the ethical implications of using LLMs, including data privacy, algorithmic bias, and potential misuse. Addressing data limitations, improving model interpretability, and establishing robust ethical guidelines are crucial for responsible innovation. The review concludes by outlining promising research directions, such as developing more sophisticated generative models, enhancing property prediction accuracy, and exploring LLM applications beyond molecule generation to unlock the full potential of LLMs in revolutionizing drug research.

Introduction

Large language models (LLMs) are rapidly transforming various scientific disciplines, and drug research is no exception. The integration of these powerful models holds the promise of accelerating the traditionally lengthy and expensive drug discovery process by facilitating tasks such as target identification, molecule generation, property prediction, and clinical trial optimization. However, the application of LLMs in this domain is still relatively nascent, demanding careful consideration of the methodologies employed, the metrics used to evaluate their performance, and the ethical implications that arise from their use. This review addresses the crucial intersection of LLMs and drug research, focusing specifically on novel methods developed to leverage LLMs for drug discovery, the evaluation metrics employed to assess the efficacy and reliability of these methods, and the ethical considerations that must be addressed to ensure responsible innovation in this field.

This review synthesizes findings from studies exploring the application of LLMs, including investigations into generative models for molecular design [1]. Specifically, we examine how these models are being adapted and refined for the unique challenges of drug discovery, moving beyond generic LLM applications. A key area of focus is the development of standardized benchmarking platforms, which facilitate the fair comparison of different generative architectures [1]. We also delve into the evaluation metrics used to assess the generated molecules, considering factors such as validity, novelty, and drug-likeness. Furthermore, this review investigates the ethical implications of using LLMs in drug research, including issues related to data privacy, bias in algorithms, and the potential for misuse.

The use of appropriate evaluation metrics and standardized benchmarks are crucial in the context of generative models such as CharRNN, VAE, AAE, JTN-VAE, and LatentGAN [1]. Each of these models requires careful scrutiny of its performance regarding its ability to generate valid, unique, and novel molecular structures while maintaining specific chemical properties. Without such standardized evaluations, a clear comparison of the strengths and weaknesses of the various architectures becomes exceedingly difficult [1]. Therefore, this review aims to provide a comprehensive overview of the current landscape, highlighting both the potential and the challenges of using LLMs to revolutionize drug research. By identifying knowledge gaps and outlining best practices, this review aims to guide future research and promote the responsible development of LLM-driven drug discovery tools. The exploration of these methods and ethics are vital to harness the full potential of LLMs for drug discovery while mitigating potential risks.

Key Themes and Trends

One key theme emerging from the research is the emphasis on standardized evaluation and benchmarking of LLMs in the context of molecular design and drug discovery. The lack of consistent evaluation metrics is a significant challenge in the field, hindering the ability to fairly compare and contrast different LLM architectures and methodologies. Specifically, the development and application of standardized platforms like MOSES are aimed directly at addressing this issue [1]. MOSES provides a curated dataset, standardized implementations of various generative models, and a suite of evaluation metrics to assess the performance of these models in generating valid, unique, and novel molecular structures while maintaining specific chemical properties [1]. This allows for a more objective comparison of different generative architectures, such as CharRNN, VAE, AAE, JTN-VAE, and LatentGAN, as highlighted by their inclusion in the MOSES benchmark [1]. The significance of this trend lies in its potential to accelerate progress in the field by providing a common ground for comparing and improving LLM-based drug discovery methods.

Another significant trend is the exploration of various LLM architectures and training techniques for molecular generation and property prediction. The selection of appropriate architectures and optimization of training parameters are crucial for achieving optimal performance. For instance, MOSES benchmark includes several models, such as CharRNN that employs GRU cells [1]. The choice of specific architectural components, like GRU cells in CharRNN, demonstrates the importance of tailoring LLM designs to the specific characteristics of molecular data. Furthermore, the optimization of hyperparameters and the use of techniques like teacher forcing [1] during training are important aspects of model development.

Furthermore, the reviewed papers demonstrate a common interest in evaluating the trade-offs between different aspects of molecular generation, such as validity, novelty, and diversity. Generating molecules that are synthetically accessible and have desirable properties is a complex task, and LLMs often need to balance these competing objectives. The MOSES platform aims to reveal the strengths and weaknesses of different generative architectures in exploration and exploitation of the chemical space [1]. This emphasizes the importance of considering multiple evaluation metrics and understanding the trade-offs between different design objectives.

Methods and Technologies

The methods and technologies employed in LLM-driven drug research encompass a diverse range of approaches, primarily focused on generative models, property prediction, and optimization techniques. A significant emphasis is placed on benchmarking and standardization to ensure fair comparisons between different methodologies [1]. Generative models, particularly, are frequently utilized for de novo molecular design, aiming to create novel compounds with desired characteristics.

One prominent method, highlighted by the introduction of the MOSES platform, involves training various state-of-the-art molecular generation models on a curated dataset and evaluating their performance using predefined metrics [1]. The generative models assessed in that paper include CharRNN (Character-level Recurrent Neural Network), VAE (Variational Autoencoder), AAE (Adversarial Autoencoder), JTN-VAE (Junction Tree Variational Autoencoder), and LatentGAN (Latent Generative Adversarial Network), alongside baseline models like HMM, N-Gram, and a combinatorial approach [1]. Each deep learning model was trained using architecture-specific hyperparameter optimization, but on the same training data [1]. CharRNN, for example, uses GRU cells in a three-layer architecture with 512-dimensional hidden states [1]. This approach allows for a direct comparison of the strengths and weaknesses of different architectures under controlled conditions.

The MOSES platform [1] is designed to standardize evaluation practices, which is a critical aspect of advancing LLM-driven drug discovery. The platform incorporates a suite of evaluation metrics designed to assess the generated molecules based on validity (whether the generated molecule is chemically valid), uniqueness (whether the generated molecule is already present in the training set), novelty (whether the

generated molecule has not been seen before), and specific chemical properties [1]. The choice of metrics is crucial, as it directly influences the perceived performance of the models. Inconsistent evaluation practices can lead to misleading conclusions about the effectiveness of different approaches, highlighting the importance of standardized platforms like MOSES [1].

Comparing different generative models reveals that certain architectures are better suited for specific tasks. For instance, junction-based models, like JTN-VAE, might excel at generating molecules with complex ring systems, whereas simpler models, like CharRNN, might be more effective at exploring a broader chemical space [1]. Understanding these trade-offs is essential for selecting the appropriate model for a given drug discovery objective. The evaluation metrics in MOSES enable researchers to quantify these trade-offs and make informed decisions about model selection [1].

Resources and Datasets

The availability and characteristics of datasets are crucial for training and evaluating LLMs in drug research. The MOSES benchmarking platform [1] provides a curated dataset specifically designed for evaluating molecular generative models. This dataset is an essential resource, enabling standardized comparisons of different generative architectures [1]. The MOSES platform also includes standardized implementations of various generative models, which facilitates their use and comparison [1]. Without such curated and standardized datasets, the comparison of different LLM approaches becomes significantly more challenging, hindering progress in the field. The curated dataset is a critical component for achieving the goal of consistent and fair comparisons among various generative architectures, highlighting their respective strengths and weaknesses [1]. The platform encompasses a suite of evaluation metrics, which are integral to assessing the validity, uniqueness, and novelty of generated molecular structures, while also considering the preservation of specific chemical properties [1]. By providing a standardized benchmark, the MOSES platform effectively establishes a foundation for future advancements in AI-driven drug discovery [1].

The MOSES platform incorporates several state-of-the-art molecular generation models, including CharRNN, VAE, AAE, JTN-VAE, and LatentGAN, alongside baseline models like HMM, N-Gram, and a combinatorial approach [1]. These models serve as important resources for researchers seeking to explore and compare different LLM architectures for molecular design. Each deep learning model within the MOSES platform was trained using architecture-specific hyperparameter optimization, yet all models were trained on the same training data, further enhancing the comparability of their results [1]. The detailed architectural specifications of models like CharRNN (using GRU cells in a three-layer architecture with 512-dimensional hidden states) provided within the MOSES documentation are also invaluable resources for replication and further development [1]. These details allow researchers to understand the specific configurations used in the benchmark and potentially adapt them for their own research.

Challenges and Limitations

A significant challenge in leveraging LLMs for drug research stems from data-related limitations. The quantity and quality of available datasets directly impact the performance and reliability of these models [1]. While resources like the MOSES platform offer curated datasets for molecular generation, the breadth and depth of chemical space remain vast, and current datasets may not adequately represent the diversity required for robust model training. Data scarcity is particularly acute for specific target classes or diseases, potentially leading to biased or poorly generalized models. Furthermore, the inherent complexity of biological data, with its multi-faceted interactions and noisy measurements, presents a challenge for LLMs to accurately capture and interpret relevant patterns. Data curation and preprocessing become critical bottlenecks, demanding significant effort and expertise to ensure data quality and relevance.

Methodological limitations also pose significant hurdles. The lack of standardized evaluation metrics and benchmarking protocols makes it difficult to compare the performance of different LLM architectures and methodologies objectively [1]. The MOSES platform attempts to address this, but the field requires broader adoption of standardized benchmarks and a more comprehensive suite of evaluation metrics that capture diverse aspects of molecular design and drug discovery. For example, beyond metrics like validity, uniqueness, and novelty, there is a need for metrics that assess the synthesizability, druglikeness, and target affinity of generated molecules. The trade-offs between exploration and exploitation in chemical space also remain a challenge [1]; while LLMs can generate novel structures, ensuring that these structures possess desirable properties and are synthetically accessible requires careful optimization and validation. Furthermore, current methods may struggle to incorporate complex biological constraints and domain knowledge effectively.

Computational constraints also factor into the challenges of applying LLMs in drug research. Training and deploying large LLMs require substantial computational resources, including high-performance computing infrastructure and specialized hardware such as GPUs or TPUs. This can limit accessibility for researchers with limited resources or those working in smaller institutions. The computational cost of generating and screening large libraries of molecules using LLMs can also be prohibitive, necessitating the development of more efficient algorithms and optimization techniques. Model interpretability is another concern, as the black-box nature of many LLMs makes it difficult to understand why a particular model generates a specific output. This lack of transparency can hinder trust and acceptance of LLM-generated predictions, particularly in high-stakes applications such as drug discovery.

Ethical considerations are increasingly important as LLMs become more integrated into drug research. Data privacy and security are paramount, particularly when dealing with sensitive patient data or proprietary chemical information. There is a risk of unintended biases being encoded in LLMs due to biases present in the training data. These biases can lead to discriminatory outcomes or the development of drugs that are only effective for certain populations. It is crucial to develop strategies for detecting and mitigating biases in LLMs to ensure fairness and equity in drug discovery. Additionally, the potential for misuse of LLMs, such as generating toxic molecules or designing drugs for nefarious purposes, raises ethical concerns that need to be addressed through responsible development and deployment practices.

Recurring obstacles across the reviewed papers include the need for better datasets, standardized evaluation metrics, and improved model interpretability. The limited availability of high-quality, diverse datasets remains a significant bottleneck [1]. The lack of standardized evaluation makes it difficult to objectively compare different models and track progress in the field [1]. The black-box nature of LLMs hinders trust and acceptance of their predictions. The implications of these limitations for the field are that progress in LLM-driven drug research may be slower than anticipated, and the potential benefits of these models may not be fully realized. Overcoming these challenges requires a concerted effort from researchers, data scientists, and policymakers to develop better datasets, evaluation metrics, and ethical guidelines for the use of LLMs in drug research.

Proposed Research Pipeline

The overarching goal of the proposed research pipeline is to streamline the application of LLMs in drug research, specifically targeting de novo molecule generation with desired properties. This pipeline aims to address the limitations of current approaches, particularly the lack of standardized evaluation metrics and the challenges associated with data scarcity and bias, as highlighted in the reviewed articles. It emphasizes a modular design allowing for iterative refinement and integration of novel methods as they emerge. The pipeline leverages insights from benchmarking platforms like MOSES [1] to ensure rigorous evaluation and comparison of different LLM architectures.

Stage 1: Data Curation and Preprocessing. This stage focuses on assembling and preparing a high-quality dataset for training and evaluating the LLM. Drawing inspiration from the MOSES platform [1], this stage will involve curating a diverse collection of chemical structures and their corresponding properties from sources like

ChEMBL, PubChem, and ZINC. Data preprocessing will include standardization of chemical representations (SMILES, InChI), removal of duplicates, and handling of erroneous or incomplete data. We will also apply data augmentation techniques, like SMILES enumeration, to expand the dataset, which mitigates data scarcity. Validation at this stage will involve checking the chemical validity and consistency of the curated data using cheminformatics toolkits like RDKit.

Stage 2: LLM Architecture Selection and Adaptation. The purpose of this stage is to select an appropriate LLM architecture and adapt it for molecular generation. Based on the benchmarking results presented in [1], we will explore different LLM architectures, including transformer-based models and recurrent neural networks (RNNs), with a focus on models demonstrated to perform well in generative tasks. The models will be adapted for SMILES string generation or graph-based molecular representations. This stage will involve modifying the model architecture and training objectives to suit the specific requirements of molecular design. Intermediate evaluation will include perplexity measurements on a held-out validation set and assessments of the model's ability to generate syntactically correct SMILES strings.

Stage 3: Property Prediction Module Integration. This stage integrates a property prediction module into the LLM-based generative framework. This module aims to predict relevant drug-like properties, such as solubility, binding affinity, and toxicity. It leverages existing QSAR/QSPR models or trains new models on the curated dataset from Stage 1, alongside other relevant datasets. The property prediction module can use machine learning algorithms such as Random Forests or Gradient Boosting, commonly used for property prediction. The LLM can then be conditioned to generate molecules with optimized predicted properties. Evaluation will involve assessing the accuracy of the property prediction module using metrics like RMSE and R-squared on a held-out test set.

Stage 4: Constrained Molecular Generation. This stage implements constrained molecular generation using the LLM and property prediction module from the previous stages. This involves guiding the LLM to generate molecules that satisfy specific property constraints or structural requirements. This can be achieved using techniques like reinforcement learning, where the LLM is rewarded for generating molecules with desired properties, or by directly incorporating property constraints into the LLM's training objective. Inspired by the work in [1] demonstrating the balance between novelty and validity, this stage will emphasize generating novel structures that also adhere to medicinal chemistry principles. Validation involves assessing the percentage of generated molecules that satisfy the specified constraints and evaluating the novelty and diversity of the generated molecules.

Stage 5: Scaffold Decoration and Optimization. This stage focuses on refining the generated molecules by modifying existing scaffolds or decorating them with new functional groups to further optimize their properties. This can be achieved using rule-based approaches, expert systems, or machine learning models trained to suggest optimal modifications. Techniques like retrosynthetic analysis can be used to identify feasible synthetic routes for the generated molecules. The goal is to enhance the drug-likeness and synthetic accessibility of the generated compounds. Evaluation involves assessing the synthetic accessibility of the generated molecules using metrics like SA score and assessing their drug-likeness using metrics like QED.

Stage 6: In Silico Validation and Filtering. This stage conducts in silico validation of the generated molecules to filter out compounds with undesirable properties or potential liabilities. This involves using computational tools to predict ADMET (Absorption, Distribution, Metabolism, Excretion, and Toxicity) properties and assess the potential for off-target effects. Molecular docking simulations can be performed to evaluate the binding affinity of the generated molecules to target proteins. Molecules that fail to meet predefined criteria are eliminated from further consideration.

Stage 7: Experimental Validation and Iterative Refinement. This stage involves experimental validation of a subset of the generated molecules to confirm their activity and properties. This can involve synthesizing the selected molecules and testing them in vitro or in vivo. The experimental results are then used to refine the LLM and the property prediction module, creating a feedback loop for continuous improvement. This stage is crucial for validating the predictions made by the LLM and ensuring the accuracy of the generated molecules.

A key challenge lies in integrating diverse data sources and ensuring data quality across the pipeline. Data harmonization and standardization are crucial for seamless data flow. Another challenge is the computational cost associated with training and deploying large LLMs. Mitigation strategies include leveraging cloud computing resources and optimizing the LLM architecture for efficient inference. Moreover, ensuring the chemical validity and synthesizability of the generated molecules remains a challenge. Integrating retrosynthetic analysis tools and medicinal chemistry rules into the generation process can help address this issue. The limited availability of experimental validation data for many compounds poses a significant hurdle. Addressing this would require prioritizing the synthesis and testing of a diverse set of molecules generated by the pipeline. Furthermore, potential biases in training data could lead to skewed results. Careful attention should be paid to data sampling and bias mitigation techniques during the data curation phase, drawing from best practices in fairness and accountability.

In summary, the proposed research pipeline provides a comprehensive framework for applying LLMs to de novo molecular design. By integrating data curation, LLM adaptation, property prediction, constrained generation, scaffold optimization, in silico validation, and experimental feedback, this pipeline aims to accelerate the discovery of novel drug candidates. The pipeline's modular design allows for continuous improvement and integration of new methods, paving the way for future advancements in AI-driven drug discovery.

Conclusion

In conclusion, the application of LLMs in drug research, as revealed by the reviewed articles, is a field brimming with potential yet facing significant hurdles. The key themes center around generative modeling, property prediction, and optimization within the drug discovery pipeline, with a strong emphasis on the crucial need for standardized evaluation metrics and benchmarking [1]. Dominant methods include the utilization of deep generative models like CharRNN, VAE, AAE, and JTN-VAE, employed for de novo molecule generation [1]. However, a major challenge lies in data limitations, encompassing both the quantity and quality of available datasets, impacting the robustness and generalizability of these models [1].

Based on the reviewed articles, the current state of research can be characterized as promising but still in its early stages. While LLMs demonstrate the capacity to generate novel molecular structures and predict drug-like properties, their performance is heavily influenced by the data they are trained on and the evaluation metrics used to assess their efficacy [1]. The MOSES platform exemplifies a concerted effort to standardize benchmarking, yet broader adoption and further refinement of evaluation protocols are warranted [1]. Without such standardization, comparing and contrasting different LLM architectures and methodologies remains problematic.

Effective research approaches should prioritize modular pipeline designs, allowing for iterative refinement and the integration of diverse data sources and evaluation metrics. A critical element is incorporating domain-specific knowledge and constraints into the model training process to enhance the relevance and applicability of the generated molecules. Moreover, the systematic exploration of different LLM architectures and training strategies is essential to identify the most suitable approaches for specific drug discovery tasks.

Promising directions for future research include developing more sophisticated generative models capable of capturing complex chemical relationships, improving the accuracy and reliability of property prediction models, and addressing the challenges of data scarcity and bias through techniques like transfer learning and data augmentation. Further investigation into incorporating structural information and biological context into LLMs is also warranted. Exploring the use of LLMs for tasks beyond molecule generation, such as predicting drug-target interactions or optimizing clinical trial designs, represents another avenue for impactful research. Addressing the ethical considerations surrounding the use of AI in drug discovery, such as data privacy and algorithmic bias, is also of paramount importance.

Advancing research in this area holds the potential to significantly accelerate the drug discovery process, reduce development costs, and ultimately lead to the discovery of novel and effective treatments for a wide range

of diseases. By addressing the identified challenges and pursuing the suggested research directions, LLMs can become a powerful tool in the fight against disease and improve human health. The development of robust, reliable, and ethically sound LLM-driven drug discovery pipelines promises a new era of pharmaceutical innovation and personalized medicine.

References

1. Adarsh Singh (2025). A Comprehensive Benchmarking Platform for Deep Generative Models in Molecular Design. arXiv preprint arXiv:2505.12848v1.